



### Domain: GCP Observability

Objective	<p>Implement a comprehensive Observability solution for GCP platform and workloads.</p> <p>As a cloud engineer, operations, or site reliability engineer, I want to implement observability tools for visibility across our GCP infrastructure to ensure that I can monitor, troubleshoot, and optimize the performance and health of my platform and application workloads.</p> <p><b>Consideration:</b></p> <ul style="list-style-type: none"><li>▪ Ability to collect and display metrics and logs for all GCP hosted applications, infrastructure, and services.<ul style="list-style-type: none"><li>○ This has a dependency on various service enablement teams for various GCP services and infrastructure.</li></ul></li><li>▪ Capability to build customized dashboards with real-time visualizations for key performance indicators for Cloud Operations and Mission Control.</li><li>▪ Support for configuration of alerts based on thresholds for key SLOs and integrate alert pipeline with Citi infrastructure and service management system.</li><li>▪ Support for end-to-end tracing of requests across application and infrastructure boundaries.</li><li>▪ Ability to correlate events from different services for faster issue resolution and root cause analysis</li></ul>
Acceptance Criteria	<ul style="list-style-type: none"><li>▪ Logs from all GCP services are collected in a central observability solution.</li><li>▪ I can search, filter, and analyze logs based on severity, labels, and time ranges.</li><li>▪ Metrics from all GCP resources are available and displayed on custom dashboards.</li><li>▪ Alerts are triggered when critical thresholds are crossed (e.g., CPU &gt; 80%, memory usage &gt; 70%).</li><li>▪ Alerts are sent to incident management tools like ServiceNow or PagerDuty.</li><li>▪ Traces are automatically collected for requests across multiple services.</li><li>▪ Latency and performance bottlenecks can be visualized in the Cloud Trace console.</li><li>▪ I can drill down into specific traces to identify the service or request causing delays.</li><li>▪ Cloud Error Reporting aggregates similar errors from application logs.</li><li>▪ Notifications are sent when a threshold of recurring errors is detected.</li><li>▪ Errors are classified by service and severity for prioritization.</li><li>▪ Uptime checks are configured for all critical services and APIs.</li><li>▪ Alerts are generated when services become unavailable or exceed latency thresholds.</li><li>▪ Reports on service uptime and availability are generated and viewable in the Cloud Monitoring dashboard</li><li>▪ SLOs are defined based on key metrics like uptime, response times, and error rates.</li><li>▪ SLO dashboards are set up, showing real-time and historical performance against targets.</li><li>▪ Alerts are sent when SLOs are at risk of being breached</li><li>▪ Retention policies should be set for different types of logs (e.g., error logs are retained for 30 days, debug logs for 7 days).</li><li>▪ Cost reports for Cloud Logging and Cloud Monitoring are generated.</li><li>▪ Log storage is optimized by exporting long-term logs to Cloud Storage.</li></ul>

Commented [F1]: Love it. This is great. I am good with this as a template. Thanks @Alsalam, Mohamed [TECH].



	<ul style="list-style-type: none"><li>▪ All resource and service limits are accessible from single view for all GCP services across different projects.</li><li>▪ I can get notification of an impending service limit breach.</li><li>▪ On-premises infrastructure is integrated into Google Cloud Monitoring and Logging.</li><li>▪ A unified dashboard displays metrics, logs, and performance data from both GCP and on-prem environments.</li><li>▪ Alerts are set for both cloud and on-prem workloads, enabling centralized incident management.</li><li>▪ VPC flow data is integrated into Google Cloud Monitoring and Logging.</li><li>▪ A unified dashboard displaying network traffic data for GCP VPCs across all projects.</li></ul>
Stakeholders	Cloud Operations, Cloud SRE, EAP, XCS and VDI team.
Resourcing	4 GCP Cloud Engineers
Milestones	<p>Dev Milestone 1 – Dec 2024</p> <ul style="list-style-type: none"><li>• Completion of GCP alert pipeline</li><li>• Integration with Netcool &amp; ServiceNow</li><li>• Certification of Cloud Trace &amp; Cloud Error Reporting</li></ul> <p>Dev Milestone 2 – Feb 2025</p> <ul style="list-style-type: none"><li>• Centralized logging across all GCP projects per environment per initiative</li><li>• Design and deployment of metrics scopes for all GCP projects per environment per initiative.</li><li>• Integration of Cloud Trace &amp; Cloud Error reporting into project vending process</li></ul> <p>Dev Milestone 3 – Mar 2025</p> <ul style="list-style-type: none"><li>• Unified Observability for dashboards for GKE based workloads (minus traces).</li><li>• Unified Observability for dashboards for GCE based workloads (minus traces).</li></ul> <p>PROD – Apr 2025</p> <p>Post PROD</p> <ul style="list-style-type: none"><li>• Service and component event tracing across all certified GCP services.</li><li>• Self Service custom resources to support service monitors</li></ul>
Submitter	Augustine Opoku.

#### For Reviewers Use Only

Feedback	
Status	